

## **Beginner Level: Setting Up the Foundation**

### 1. Kubernetes Basics

- Objective: Set up a basic Kubernetes cluster on your local machine.
- Steps:
  - Install Minikube or Kind (Kubernetes in Docker) for local Kubernetes clusters.
  - Get familiar with `kubectl` commands (e.g., `kubectl get`, `kubectl describe`, `kubectl apply`).
  - Create simple YAML files for Pods, Services, and Deployments.
- Outcome: Successfully deploy a single-container application (e.g., Nginx or a simple Node.js app) on the Kubernetes cluster.

### 2. Basic Service Deployment

- Objective: A microservice based application is provided to you.  
Project: <https://github.com/UnpredictablePrashant/Slab.ai.git>
- Steps:
  - Write a simple REST API for listing products (e.g., using Node.js or Python).
  - Create Docker images for the application.
  - Expose the service with a LoadBalancer or NodePort Service.
- Outcome: Access the Product Catalog service via a public URL or IP.

## **Intermediate Level: Scaling and Networking Microservices**

### 3. Creating Additional Microservices and Networking

- Objective: Add more services and enable internal communication.
- Steps:
  - Create two additional microservices: User Authentication and Order Management.
  - Use Kubernetes Services to enable these microservices to communicate.
  - Configure Ingress resources for external access and route traffic based on paths.
- Outcome: Deploy an Ingress controller (e.g., Nginx Ingress) and configure rules to route traffic to different microservices.

### 4. ConfigMaps and Secrets

- Objective: Manage configuration and sensitive data.
- Steps:
  - Use ConfigMaps to manage environment variables and application configuration.
  - Use Secrets to securely store sensitive data like API keys and passwords.
  - Update your services to consume these ConfigMaps and Secrets.
- Outcome: Configurations and sensitive data are stored securely, and services use them dynamically.

### 5. Auto-scaling and Load Management

- Objective: Implement horizontal pod autoscaling based on resource utilization.
- Steps:
  - Configure resource limits for each service.
  - Set up Horizontal Pod Autoscalers (HPAs) to scale based on CPU utilization.
- Outcome: Kubernetes scales services automatically based on load, ensuring efficient resource usage.

### **Advanced Level: Production-Ready Enhancements**

#### 6. Stateful Applications and Storage

- Objective: Deploy a Database (e.g., MongoDB or MySQL) and manage data persistence.
- Steps:
  - Deploy a database using StatefulSets and PersistentVolumeClaims for data persistence.
  - Configure services to connect to the database.
  - Implement a backup strategy for your database data.
- Outcome: Your database persists data even when Pods are deleted or restarted.

#### 7. CI/CD Integration with GitOps

- Objective: Implement a CI/CD pipeline to automate deployments.
- Steps:
  - Set up a Git repository to store your Kubernetes manifests.
  - Use a tool like ArgoCD or Flux for GitOps-based deployment.
  - Configure a CI/CD pipeline in Jenkins, GitHub Actions, or GitLab CI/CD to automate builds and deployments to the Kubernetes cluster.
- Outcome: Every change to the repository automatically deploys to the Kubernetes cluster.

#### 8. Monitoring and Logging

- Objective: Implement monitoring and logging for application and cluster insights.
- Steps:
  - Deploy Prometheus and Grafana for monitoring, setting up dashboards for CPU, memory, and request metrics.
  - Integrate a logging solution (e.g., EFK stack - Elasticsearch, Fluentd, and Kibana).
  - Configure alerts in Prometheus based on thresholds.
- Outcome: Gain insights into the performance and health of your services and the Kubernetes cluster.

#### 9. Service Mesh Integration

- Objective: Add a service mesh like Istio or Linkerd for advanced networking features.
- Steps:
  - Install a service mesh and configure it to manage traffic between microservices.
  - Enable features like mutual TLS for security, traffic splitting, and circuit-breaking for resilience.
- Observe telemetry and distributed tracing for inter-service calls.

- Outcome: Enhanced microservices communication with security and resilience.

## Project Wrap-Up: Deploying to a Managed Kubernetes Cluster

### 10. Deploying to a Cloud Provider

- Objective: Move your local Kubernetes setup to a managed cloud Kubernetes service.
- Steps:
  - Choose a cloud provider (AWS EKS, Google GKE, or Azure AKS) and create a Kubernetes cluster.
  - Use kubectl to deploy all your configurations to the cloud cluster.
  - Configure cloud-specific features like managed storage and load balancers.
- Outcome: Your application is running on a managed Kubernetes cluster, ready for production.